

AA-10

Manager Strategic Roadmap

Ten weeks, four engineers, the Gantt, the risk register and the ordered cut list

30 August 2026 · Rev 1.0 — baseline · AA-10 of thirteen

ENGINEERS ASKED FOR

4

CALENDAR

10 weeks

DEV-DAYS COMMITTED

200

RECOVERABLE BY CUTS

20 days

RISKS TRACKED

12

RECURRING COST

Rs 0

This is the document for the person who has to commit the engineers and then answer for that commitment every Friday. It does not explain how anything is built; AA-00 gives the architecture and AA-11 gives the task-level plan. It states what is being asked for, what arrives at the end of ten weeks, what will not arrive, the decisions you personally have to make and by which week, and the three conditions under which stopping this programme is the correct call rather than a failure.

10.1 THE COMMITMENT IN ONE PAGE

Four engineers, ten weeks, one hire that does not exist yet

ENGINEERS ASKED FOR

4

BE-1, BE-2, FE-1 and a new FE-2

CALENDAR

10 weeks

W1 to W10, pilot inside W10

DEV-DAYS COMMITTED

200

20 of them recoverable by cuts

PORTAL SCREENS

91

41 admin, 26 supplier, 24 fleet

NEW TABLES

36

8 migrations, M1 to M8

RECURRING COST

Rs 0

the ₹67 domain is already paid

What is being asked for is 4 engineers for 10 consecutive weeks: 200 dev-days, of which 20 are pre-agreed as cuttable and the rest are not. Two of the four are the existing backend engineers. The third is the existing front-end engineer, who spends week one building a shared component kit rather than screens. The fourth, FE-2, does not exist yet — the plan assumes a front-end hire on the tools in week one and building supplier screens by week two. MOB-1 is deliberately not part of this commitment; the mobile app is touched exactly once, by a single order-tracker change in week four, and otherwise continues on its own roadmap.

What arrives at the end of week ten is three browser applications and the platform changes behind them. An admin portal of 41 screens replaces the founder console and adds the agronomist queue. A supplier portal of 26 screens lets a product company onboard itself, submit KYC, draft products into moderation, and then accept, pack and hand over its own shipments. A fleet portal of 24 screens lets a third-party delivery company accept broadcast offers, attach riders, and complete deliveries with an OTP inside a 200 m geofence. Underneath, orders split into per-supplier shipments so that two suppliers on one order stop being a manual reconciliation problem, and both partner commission and fleet delivery fees accrue into append-only ledgers that settle on a weekly payout run.

The business case is one sentence. The platform today cannot accept a new supplier or a new delivery partner without an Atavishaala employee doing the onboarding, the cataloguing, the shipment coordination and the payout arithmetic by hand — and that employee is the ceiling on how many suppliers and how many deliveries the business can carry. This programme removes that ceiling and replaces it with a permission matrix and an audit log. It is an operating-model change delivered as software, not a feature release.

What does not arrive matters as much, because a manager who assumes any of the following is included will be surprised in week ten. The table below lists the seven things most likely to be assumed. Note the difference between the first six, which are deferred, and wallet cash-out, which is not deferred but prohibited: paying money out of a stored balance needs an RBI prepaid-instrument licence, so it is not a scheduling question at all.

NOT DELIVERED	WHY NOT	WHAT HAPPENS INSTEAD
A native rider application	A rider uses the fleet portal in a mobile browser. There is no app store build in this plan.	Riders need a working phone browser and a data connection at the doorstep. If that proves unworkable in the pilot, a rider app is a Phase 2 programme of its own, not a cut.
Live GPS tracking on a map	Continuous location ingest and map tiles are both billed, and neither fits Rs 0.	The customer sees milestone statuses only. The POD captures one geo point, which is all the 200 m geofence check needs.
Masked calling between rider and customer	No free provider offers a number-masking call bridge in India.	The rider gets the address on a job-scoped manifest and contacts the customer through a WhatsApp template. The customer phone number is never released to a fleet.
GST-compliant invoices generated by the platform	Tax document generation is a compliance surface, not a sprint task.	The self-billed commission invoice is on the cut list as C-3. Accounting issues invoices manually for the pilot suppliers.
A public partner or fleet API	No partner has asked with a system to integrate. Building it first would be speculative.	The webhook_deliveries table exists from M8 and stays empty until someone needs it.
Custom, customer-defined staff roles	Nine system roles cover every seat that can be named today.	On the cut list as C-4. If it survives the ten weeks it ships; if week eight is late it is the second cut taken.
Wallet cash-out to a bank account	Prohibited, not deferred. Paying out of a stored value balance requires an RBI prepaid-payment-instrument licence.	Wallet credit is spendable inside Atavishaala and refundable to the original payment instrument. It never becomes cash. This does not change in Phase 2.

WHERE THE DETAIL LIVES, SO THIS DOCUMENT CAN STAY SHORT

AA-00 sections 0.11 to 0.16 hold the charter view of the same plan: the team table, five headline risks, the Rs 0 budget position, the invariants and the definition of done. This document does not repeat them; it extends them with the full twelve-risk register, the ordered cut list, the Gantt and the decision calendar. AA-01 has the five-repository split and why the portals are separate deployments. AA-09 has the per-service free-tier arithmetic and every alert threshold behind the Rs 0 claim. AA-11 has all 130 tasks with owners and acceptance criteria, which is the document the engineers work from. AA-12 is the two-page summary to hand upward; it is deliberately derived from this one so the numbers cannot diverge.

10.2 THE TEN WEEKS

Read-only before write, tenancy before features, money last

WK	THEME	DEMONSTRABLE END STATE	WHAT TO WATCH THAT WEEK
W1	Foundations and identity	Five repos exist, CI green on all five, staff login with TOTP works end to end, the kit publishes its first version.	The kit is the bottleneck. If FE-1 has not published v0.1 by Wednesday, FE-2 is idle.
W2	Permissions, audit and the admin shell	The 88 permissions are seeded and enforced; every admin mutation is audited; the admin shell renders 12 read-only screens against real data.	DENY-wins resolution and the masking extension must both land before any list screen.
W3	Partner tenancy and the moderation pipeline	A product company can be invited, submit KYC, be verified, draft a product and submit it; a catalog manager can approve it and it appears in the app.	First end-to-end cross-portal flow. Demo it live at the Friday gate.
W4	The shipment split	Orders split into shipments per supplier; orders.status is derived; the app tracker still reads correctly; M4 is proven on a restored snapshot.	The single riskiest week. Nothing else is scheduled in parallel for BE-2.
W5	Offer allocation and the fleet marketplace	A ready shipment broadcasts an offer; a delivery company accepts it, attaches a rider and a vehicle; unaccepted jobs escalate through three waves onto the admin board.	Allocation fairness and the expiry sweep are the two things to hammer in testing.
W6	Delivery execution, proof of delivery and disclosure	A rider scans pickup, opens the manifest (logged), delivers with OTP inside the geofence, and the order shows DELIVERED in the app. Failed attempts and RTO both work.	The PII disclosure register must be populating correctly before this week closes.
W7	Settlement and cash on delivery	Payout accrual, weekly run assembly, lock, release with UTR capture; COD declaration and reconciliation; a synthetic week reconciled to the paisa.	Finance-facing. Nothing here may be approximate.
W8	Returns, RTO and the Advisor console	The full return and RTO loop with compensating ledger entries, and agronomists working the advisory queue from the admin portal within the business-hours SLA.	Two unrelated workstreams; keep them on separate branches to avoid a tangled merge.
W9	Hardening, contract sweep and the free-tier audit	Twelve CI gates green on five repos, the PII sweep clean, rate limits and CSP in place, eight free-tier thresholds configured and proven to fire.	Hardening is not optional polish; the PII sweep is a gate condition.
W10	Pilot, runbooks and cutover	Two real product companies and one real delivery company transact end to end. Eleven runbooks written and rehearsed. Go or no-go decided on evidence.	No new features. Bug budget only. The cut list is exercised here if needed.

The order of the weeks is not the order the screens are numbered in, and it is not the order an enthusiastic team would pick. Three rules produced it. First, read-only before write: weeks one and two ship identity, the permission matrix, the audit middleware, the masking extension and twelve read-only admin screens, and no mutation reaches production until the thing that records mutations is already there. The rejected alternative was to build the moderation queue first because it is the most visible win; it lost because a moderation action taken before the audit log exists is an action nobody can reconstruct, and the moderation queue is exactly where a wrong call gets expensive.

Second, tenancy before features. Week three establishes that a partner organisation exists, has users, has a KYC state and has a scope claim in its token, before week four gives it shipments to act on and week five gives fleets offers to accept. Retrofitting tenancy is the single most reliable way to produce a data leak: every query written before the scope rule exists is a query that has to be found again and fixed. Doing it in week three means the scope filter is in the repository layer before there is anything worth filtering.

Third, money last. Settlement, COD reconciliation and payout runs are week seven, after the shipment split in week four and delivery execution in week six, because a ledger built on a fulfilment model that is still moving will be a ledger with compensating entries for engineering mistakes rather than for business events. The cost of this ordering is real: nothing about the commission economics is provable until week seven, which is late for a decision that could invalidate the fleet marketplace. Section 10.11 treats that as one of the three kill conditions

rather than pretending the sequencing is free.

The two load-bearing weeks

W1 IS LOAD-BEARING BECAUSE OF THE KIT, NOT THE CODE

The week-one deliverable that matters is not staff login, it is the first published version of the shared component kit. The kit is the bottleneck. If FE-1 has not published v0.1 by Wednesday, FE-2 is idle. A kit that slips to the following Monday costs FE-2 two or three days of the ten weeks immediately, and because there is no slack week anywhere before hardening, that loss does not get absorbed — it moves the end date by the same amount. This is the cheapest week to protect and the one most likely to be treated as setup. If the Wednesday blocker sync in week one reports the kit as not yet published, that is a red week, not an amber one.

W4 IS LOAD-BEARING BECAUSE THREE LATER WEEKS DEPEND ON IT

The shipment split is the only structural change to existing production data in the whole programme. The single riskiest week. Nothing else is scheduled in parallel for BE-2. Offer allocation in week five, delivery execution in week six and settlement in week seven all operate on shipments, so a one-week slip in week four pushes each of them by a week and lands on week nine. Weeks nine and ten cannot compress: hardening is a gate condition and the pilot has external participants with their own calendars. A week lost in week four is therefore paid for either by shipping without the PII sweep, which is not permitted, or by moving the cutover date.

The other eight weeks are recoverable, and that is a deliberate property of the plan rather than luck. Weeks two, three, five, six, seven and eight each end in a demonstrable state that can be reached with a reduced feature set, and the cut list in section 10.6 names which feature comes off in which of those weeks. Week nine can be shortened only by removing scope from earlier weeks, never by shortening hardening itself. Week ten has no engineering content at all beyond a bug budget, which is what makes it usable as the shock absorber it will probably have to be.

10.3 THE GANTT

Twenty workstreams, ten weeks, and the fortnight that is too full

WORKSTREAM	OWNER	1	2	3	4	5	6	7	8	9	10
Repo scaffolding & CI	FE-1										
Shared portal kit	FE-1										
Staff identity & TOTP	BE-1										
Permissions & audit	BE-1										
PII masking & disclosure	BE-1										
Admin portal screens	FE-1										
Partner tenancy & KYC	BE-1										
Moderation pipeline	BE-1										
Supplier portal screens	FE-2										
Shipment split (M4)	BE-2										
Fleet tenancy & areas	BE-2										
Offer allocation engine	BE-2										
Delivery execution & POD	BE-2										
Fleet portal screens	FE-2										
Settlement & payout runs	BE-2										
COD reconciliation	BE-2										
Returns & RTO	BE-2										
Advisor console	BE-2										
Hardening & contract sweep	all										
Pilot & runbooks	all										

LEVEL	FILL	WHAT IT MEANS FOR ONE ENGINEER IN ONE WEEK
3	primary	This is the engineer's week. Four to five days on this workstream and nothing else is promised. Two level-3 bars in the same column for the same owner is a planning error, not ambition.
2	substantial	Two to three days. Real progress is expected and a Friday demo will show it, but the engineer is sharing the week with something else.
1	light	Half a day to a day: review, a small fix, a contract change, or supporting somebody else who is at level 3 on the same surface.

Read the chart by column, not by row. A row tells you how long a workstream is alive, which is mostly uninteresting; a column tells you how many things one of four engineers is holding at once, which is the only thing that predicts a missed Friday gate. Week 9 has 16 of the twenty workstreams active at some level, the most of any week, and the busiest fortnight by total loading is weeks 2 and 3.

The parallelism in weeks one to four is real, because it runs along the owner boundaries in section 10.4. BE-1 is on identity, permissions and masking while BE-2 is on the shipment split, and the two touch different modules and different tables; FE-1 is on the kit and the read-only admin screens while FE-2 is consuming the kit for supplier screens. Nothing in that first block requires two people to agree on the same file, which is why four bars in one column there is not a warning sign.

The parallelism in weeks five to seven is optimistic, and the chart should be read with suspicion there. Offer allocation, delivery execution, settlement and COD reconciliation are all owned by BE-2 and all sit on the same three tables, so the bars are sequential in reality even where the chart shows them touching. That is the concrete meaning of the note in section 10.4 that BE-2 is the critical path: the only way to widen weeks five to seven is to move a workstream off BE-2 entirely, and the only candidate is returns and RTO, which BE-1 could take at the cost of the moderation polish in week eight.

Two bars are deliberately drawn where a manager might not expect them. PII masking and disclosure returns at level 3 in week six because that is when the manifest endpoint is built, and it is the single place in the whole system where an address leaves the database — it gets a full week of one engineer, not an afternoon. Hardening

reaches level 3 only in week nine, but it carries a level-1 bar in weeks two and eight, because the contract sweep and the permission-matrix test are written alongside the features they cover rather than retrofitted in the hardening week.

10.4 THE TEAM

Four engineers, five seats, and one of them is not hired

DEV	ROLE	MODULES AND REPOSITORIES	WHAT THEY OWN OUTRIGHT
BE-1	Backend — platform	backoffice, partners, catalog, identity, platform	Owns the 88-permission matrix, the masking extension, the audit middleware, the moderation pipeline and partner settlement.
BE-2	Backend — fulfilment	orders, fleet, money, advisor	Owns the shipment split, both FSMs, the offer allocator, POD validation, COD reconciliation and fleet payouts. The critical path.
FE-1	Front end — admin + kit	atavishaala-portal-kit, atavishaala-admin-portal	Owns the shared kit first, then 41 admin screens including the Advisor console.
FE-2	Front end — partner + fleet	atavishaala-supplier-portal, atavishaala-fleet-portal	26 supplier screens then 24 fleet screens. Reuses the kit; contributes back, never forks.
MOB-1	Mobile — unaffected	atavishaala-api /app namespace consumer	Continues app work. Touched only by the order-tracker change in Week 4.

Ownership here is not a suggestion about who is available. It is the boundary that makes the parallelism in section 10.3 possible, and it is drawn along module lines from AA-02 so that two engineers are rarely in the same file. BE-1 owns everything that decides whether a request is allowed — the permission matrix, the masking extension, the audit middleware — and BE-2 owns everything that moves a shipment or a rupee. When those two need to meet, they meet at the contract package, not in a merge conflict.

MOB-1 is listed and then deliberately excluded from the commitment. The mobile application is touched exactly once, by the order-tracker change in week four, and that change is designed to be additive: the tracker reads `order_events`, so splitting an order into shipments does not alter the response shape the app already consumes. If you are being asked to fund five engineers for ten weeks, the ask has grown beyond this plan.

The FE-2 hire is your decision, and the plan does not work without it

The plan assumes a second front-end engineer is on the tools in week one and building supplier screens in week two. FE-2 carries fifty supplier and fleet screens; FE-1 carries 41 admin screens and the kit that all three portals import. There is no version of this plan where FE-1 absorbs both, and pretending otherwise is how a ten-week programme becomes a sixteen-week one that nobody decided to extend.

If the hire does not land by the end of week two, the following happens in this order, and it is worth agreeing the order now rather than in week five. First, the fleet portal slips behind the supplier portal, so weeks five and six deliver the fleet API with the internal ATV-SELF fleet operated from admin screens instead of a fleet-facing portal. Second, the entire cut list is taken, all eight items, recovering 20 dev-days. Third, the supplier portal ships with the ten screens that the pilot suppliers actually need — onboarding, KYC, the product drafts, the shipment queue and the payout view — and the remaining sixteen become Phase 2.

The arithmetic is unforgiving. FE-2's share of the commitment is 50 dev-days out of 200. The whole cut list recovers 20 dev-days, and only 13 of those sit in the supplier and fleet portals FE-2 was going to build; the other 7 are admin screens FE-1 owns and cutting them frees the wrong person. So a missing FE-2 is a deficit of 50 front-end days against a recovery of 13 in the same lane, and the gap is closed by shipping fewer screens, not by working harder. R-12 in the register below is rated high probability for exactly this reason.

THE ONE THING NOT TO DO WHEN FE-2 IS LATE

Do not let FE-1 write supplier or fleet screens while also owning the kit. It is the most natural response and it is the one that costs the most, because the kit stops shipping its Wednesday release the moment its owner has screen deadlines, and every portal then waits on components that never land. R-05 and R-11 both describe what happens next: a component gets copied into a portal repo, then auth handling diverges between the supplier and fleet portals, and the programme acquires two authentication surfaces to audit instead of one. If FE-1 must take screens, take the kit off the critical path first by freezing it at the week-four release and accepting that later screens use whatever it contains.

10.5 THE RISK REGISTER

Twelve risks, an owner each, and the indicator that tells you it has started

ID	RISK	P	I	OWNER	LEADING INDICATOR / TRIGGER	RESPONSE
R-01	The M4 shipment backfill corrupts live order history	M	H	BE-2	Any mismatch between derived and stored status on the snapshot rehearsal	Rehearse on a restored Neon branch three times. Ship M4 behind a read-only window of 15 minutes. Keep the pre-migration status in orders.legacy_status for two releases.
R-02	No third-party fleet onboards in time for the pilot	H	H	You	Fewer than two fleet orgs at KYC_VERIFIED by end of W8	The ATV-SELF internal fleet org seeded in M5 is the fallback: staff deliver, the portal is used by staff, and the product is still complete. Do not delay go-live for this.
R-03	Offer allocation starves small fleets or thrashes	M	M	BE-2	Acceptance rate below 40 % or median accept latency above 8 minutes	Wave 1 is tier-1 fleets in the pincode, wave 2 adds tier 2, wave 3 is all. Rotate the within-wave order by least-recently-offered, not by score, until there is real volume.
R-04	A rider needs the customer phone and the masked number blocks delivery	H	M	BE-1	Any support ticket where a delivery failed because the rider could not call	This is a known, accepted gap. MVP answer: a time-boxed reveal on the manifest that logs to pii_disclosures. Named upgrade: an Exotel or Kaleyra click-to-call bridge at roughly ₹0.60 per minute, requiring a written budget waiver against I-14.
R-05	The kit becomes a bottleneck and FE-2 forks it	M	M	FE-1	A component copied into a portal repo instead of added to the kit	FE-2 has direct commit rights on the kit. A weekly kit release on Wednesday, and a CI check that flags a new file under portal/src/components that duplicates a kit export name.
R-06	Three portals plus the app exceed the Cloud Run always-free ceiling	L	H	BE-1	Monthly vCPU-seconds above 140k, or requests above 1.6 M	Portals are low-traffic by nature: tens of staff, tens of suppliers, hundreds of riders. The tripwire is the alert at 70 % in AA-09. Response is edge caching on the read-heavy list endpoints and raising the minimum poll interval on the offer stream.
ID	RISK	P	I	OWNER	LEADING INDICATOR / TRIGGER	RESPONSE
R-07	Neon 0.5 GB storage is exhausted by audit_log and pii_disclosures	M	M	BE-1	Database size above 350 MB, or audit_log above 2 M rows	Diff JSON is capped at 8 KB per row with truncation noted. A monthly cron exports audit rows older than 180 days to a Cloudflare R2 free-tier bucket as gzipped NDJSON, then deletes them. pii_disclosures is never purged; it is small.
R-08	A partner disputes a settlement figure and there is no evidence trail	M	H	BE-2	Any dispute that cannot be answered from partner_ledger alone	Append-only ledger plus an immutable pricing_snapshot per order line plus the effective-dated rate card. Never edit a payout line; post a compensating entry.
R-09	COD cash goes missing between rider and platform	H	H	You	Variance above 2 % of expected in any weekly period	A fleet with a negative ledger balance is blocked from new offers automatically. Weekly declaration is mandatory. Above 2 % variance, the fleet moves to prepaid-only orders until the period is reconciled.
R-10	A staff member is phished and the admin portal is used to exfiltrate data	L	H	BE-1	Any export volume anomaly, or a login from a new country	Mandatory TOTP, no raw PII anywhere so an export is of limited value, step-up on every destructive action, and an export cap of 5 000 rows per job with all exports audited.
R-11	The supplier and fleet portals ship with different auth bugs	M	M	FE-2	Any auth logic written inside a portal repo rather than the kit	Session handling, refresh rotation, the permission gate and the 401 interceptor live in the kit only. A portal that imports axios or fetch directly fails the ESLint rule.
R-12	Ten weeks with four engineers is optimistic	H	M	You	Any Friday gate missed two weeks running	The AA-10 cut list is ordered and pre-agreed. C-1 to C-4 recover 11 dev-days without touching the core loop; the whole list recovers 20, which is one week of team capacity and no more. Below the line, nothing is cuttable.
CODE	PROBABILITY (P) MEANS				IMPACT (I) MEANS	
H	More likely than not to happen during the ten weeks. Plan as though it will.				Moves the end date, or puts data or money at risk. You will hear about it.	

CODE	PROBABILITY (P) MEANS	IMPACT (I) MEANS
M	Plausible. Has happened on comparable work. Watch the indicator weekly.	Costs days and rework, absorbed by a cut rather than by a date change.
L	Unlikely, but the cost of being wrong is what earns it a row here.	Contained by a threshold or a runbook; recoverable within a week.

Twelve risks is not a long register for a programme this size, and that is deliberate: a register nobody reads is worse than none. Every row has one named owner, one indicator that can be checked at the Friday gate without asking an engineer to estimate anything, and a response that is already funded inside the plan. Ranking by probability multiplied by impact puts R-02 and R-09 at the top with the maximum score, followed by 4 risks in the next band. Two of that band are handled elsewhere in this document: R-12 is the FE-2 arithmetic in section 10.4, and R-04 is the masked-calling purchase decision in section 10.7. The three below are the ones that will actually consume your attention.

R-02 — no third-party fleet onboards in time for the pilot

This is the risk that makes the delivery-marketplace thesis either true or premature, and it is rated high on both axes because nothing in the ten weeks can make a delivery company sign. The engineering is the easy half. The indicator is countable and it is not a judgement call: fewer than two fleet organisations at KYC_VERIFIED by the end of week eight. Because fleet onboarding is manual until the portal exists, that count is visible from week five onwards, which gives you three weeks of warning rather than a surprise in week ten.

The response is already built and this is the important part: the ATV-SELF internal fleet organisation is seeded by migration M5 in week five, so the fleet portal has a tenant on day one regardless of whether anyone external signs. If no fleet onboards, Atavishaala staff operate the fleet portal themselves, deliveries still run through offers, POD, geofence and the fleet ledger, and every line of the fleet code path is exercised in production. The product is complete; only the marketplace is unproven.

The decision you will have to make is whether the pilot cutover waits for an external fleet. The recommendation in this plan is no, and the reason is that waiting converts a marketing problem into a schedule problem: the team goes idle, the code goes stale, and the commission economics stay untested either way. Ship on ATV-SELF, then onboard fleets against a working system you can demonstrate, which is a materially easier sale than onboarding them against a promise. What that costs you is honesty in reporting — a go-live with no third-party fleet must be described as such and not as a delivery marketplace launch.

R-09 — COD cash goes missing between rider and platform

Cash on delivery puts real money in the hands of people employed by a company that is not yours, and it is rated high on both axes because at pilot volumes a single bad week is a meaningful fraction of margin. The indicator is a computed number, not an impression: variance above 200 basis points of expected collection in any weekly period, which the COD reconciliation screen calculates from cod_remittances against the shipment values. The check is part of the Friday 17:00 budget-and-threshold slot, so it happens whether or not anyone remembers to worry about it.

Two controls run automatically and both are inside the ten weeks. A fleet whose ledger balance goes negative is blocked from receiving new offers with no human in the loop, which caps exposure at roughly one collection cycle rather than letting it compound. A fleet whose variance exceeds two per cent moves to prepaid-only orders until the period reconciles, which means it keeps working and keeps earning while the cash question is settled — a suspension that stops a fleet earning is a suspension that guarantees you never get the money back.

Your decision is the write-off threshold and who signs it, and it is better decided in week seven than in the week it first happens. The plan deliberately does not automate a write-off: a shortfall is recorded as a SHORT remittance and stays visible on the exception board until a named person posts a compensating ledger entry against it. That is slower than an automatic adjustment and it is the right trade, because the alternative is a system that quietly absorbs losses at a rate nobody is reporting.

R-01 — the M4 shipment backfill corrupts live order history

This is the only risk in the register that can damage data that already exists. Migration M4 splits every historical order into shipments and switches orders.status from a stored column to a projection derived by

reduceOrderStatus, and a mistake there is not a bug that gets fixed next release — it is an order history that no longer matches what customers were told. Probability is medium because the transformation is mechanical and testable; impact is high because the blast radius is the whole table.

The indicator is a zero-tolerance one and it is checked before production, not after: any mismatch between the derived status and the stored status when the migration is rehearsed on a restored Neon branch. D-18 in the definition of done requires a zero-row diff, and the rehearsal happens three times. On the day, M4 ships behind a fifteen-minute read-only window, and the pre-migration value is retained in orders.legacy_status for two releases so that the comparison remains possible after the fact rather than only before it.

There is no decision for you inside R-01, which is why it sits third rather than first, but there is one gate: the week-four Friday demo either shows a zero-row diff on a restored snapshot or it does not, and if it does not, the correct response is to hold M4 and let weeks five to seven start on a feature-flagged path. Holding M4 costs the schedule a week. Shipping a backfill that has not produced a clean diff costs credibility that the programme does not have spare.

10.6 THE CUT LIST, IN ORDER

Eight pre-agreed cuts worth 20 dev-days, and eight things that are not for sale

ID	WHAT COMES OFF	WHAT SHIPS INSTEAD	SAVING	CONSEQUENCE YOU ACCEPT
C-1	Fleet SSE offer stream	Fall back to 10 s polling everywhere.	2 dev-days	Riders will not notice; dispatch sees a few seconds of extra latency.
C-2	Saved views and CSV export	Ship fixed default filters only.	3 dev-days	Ops will ask for it within a month. It is a strong Phase 2 candidate.
C-3	Self-billed commission invoice PDF	Supplier downloads a CSV of the run instead.	2 dev-days	Accounting can generate the invoice manually for two pilot suppliers.
C-4	Custom roles	Ship the nine system roles only; no role editor screen.	4 dev-days	Nine roles cover every seat we can name today.
C-5	Rider mobile scan camera	Manual shipment-code entry only.	2 dev-days	Slower at the gate. Acceptable for a pilot with two suppliers.
C-6	POD photo capture	OTP plus geofence only; no photo.	1 dev-day	Weakens dispute evidence. Take this cut only for equipment-free, low-value routes.
C-7	Vehicle registry	Riders exist; vehicles do not. Heavy equipment goes to ATV-SELF.	3 dev-days	Blocks third-party heavy-equipment delivery. Acceptable at pilot scale.
C-8	Service-area approval workflow	Admin edits fleet areas directly.	3 dev-days	Fine with three fleets. Does not scale past ten.

A cut list is only useful if it was written before anybody was under pressure, and if the order is fixed. Both are true here. Cuts are taken in ID order, from C-1 downwards, because the list is sorted by how little the business notices: C-1 costs a few seconds of dispatch latency, C-8 costs a workflow that does not scale past ten fleets. Taking them out of order means taking a larger consequence for the same number of days, which is the mistake a team makes at 6pm on a Thursday when the nearest cut is the one they were already annoyed by.

Cuts are taken at the Friday gate and never mid-week. The reason is not ceremony: a cut taken mid-week abandons work already started, so the days it appears to save are days already spent. At the gate the week is either passed or explicitly failed, and a failed gate is the moment the next cut on the list is taken — one cut per failed gate, announced at the Monday kick-off, with the cut-list status stated even in the weeks when nothing is cut.

The whole list is worth 20 dev-days. With 4 engineers that is 5 working days of team capacity, which is one calendar week and no more. This is the number most likely to be misread in this document, so it is worth stating plainly: the cut list does not buy a month. It buys roughly one week, and it buys it unevenly — 13 of the 20 days are supplier and fleet portal work and 7 are admin work, so the relief only lands on the engineer whose cuts you took. Every cut is reversible in the days named, which is what makes the list safe to use; none of them requires undoing a migration or a data shape.

WHAT THE CUT LIST CANNOT DO

Nothing on this list moves the end date by more than about a week, and nothing on it substitutes for the FE-2 hire. If you find yourself needing three weeks, the cut list is the wrong instrument and there are only two honest levers left: move the pilot cutover date, or reduce the pilot to one supplier and the internal ATV-SELF fleet so that weeks five to eight carry fewer external dependencies. Both are decisions for you rather than for the team, which is why section 10.7 gives each of them a week by which it has to be made.

Below the line: what is not on the cut list at any price

#	NOT CUTTABLE	WHY IT IS BELOW THE LINE
1	The 88-permission matrix with DENY-wins. A portal without authorisation is not shippable.	Cutting this ships a portal where authorisation is advisory. There is no partial version: either every route resolves an effective permission with DENY winning, or the tenancy story in AA-02 is decoration.
2	The last-four-only masking extension and the pii_disclosures register.	This is the rule the whole PII regime rests on, and it is enforced by absence — the raw-read permission does not exist in the type union. Adding it later to save days is not a cut, it is a different product with a different risk profile.
3	assertTransition on shipments and delivery jobs. Without it the data becomes unauditale.	A state machine without a transition guard produces order and job histories that cannot be reconstructed, and both dispute processes in AA-09 depend on reconstructing them.
4	The append-only ledgers and the integer-paise rule.	Append-only ledgers in integer paise are what make a settlement dispute answerable from data rather than from memory. R-08 has no response at all without this.
5	Idempotency keys on money and on offer acceptance.	Without idempotency, a retried payout or a double-tapped offer acceptance becomes a real duplicate. This is two lines of middleware and one unique index; there is nothing to save.
6	Audit logging of every admin mutation.	The audit log is the only record of who did what in a portal that nine roles can reach. It is also the control that makes R-10 survivable, because a phished account leaves a trail.
7	Mandatory TOTP for staff.	One shared password is one incident away from full back-office access. TOTP is a week-one task and stays a week-one task.
8	The M4 backfill rehearsal on a restored snapshot.	The rehearsal is the entire mitigation for R-01, the only risk that can damage data that already exists. Skipping it saves two days and risks the order history.

These eight are not a wish list and they are not the team protecting its favourite work. Each one maps onto an invariant from AA-00 section 0.14 and onto at least one row of the risk register, and each one is cheap to build now and expensive to retrofit. That asymmetry is the whole argument: the masking extension is a day in week two and a full audit of every query in week twelve; TOTP is a task and a migration in week one and an incident response later. If a week is failing badly enough that one of these looks tempting, the correct action is to take the next cut on the list above and, if the list is exhausted, to move the date.

10.7 THE DECISIONS YOU HAVE TO MAKE, WITH DATES

Eight decisions that belong to you, not to the team

ID	THE DECISION	BY	WHAT HAPPENS IF IT IS DEFERRED	THE COST OF SAYING YES
K-1	Hire FE-2, or accept the reduced screen scope	W1	Every week of delay is a week of supplier and fleet screens that does not get built, and the capacity cannot be recovered later because there is no slack week before hardening. Deferred past W2, this stops being a hire decision and becomes a scope decision by default.	One front-end salary for ten weeks. The alternative is roughly forty supplier and fleet screens moving to Phase 2.
K-2	Freeze the pilot supplier and fleet list	W3	Partner tenancy, KYC document types and the rate-card shape are all built in W3. A supplier added in W6 with a different document set is a migration, not a data-entry task.	Free, but it commits you to two suppliers whose categories the catalog already supports.
K-3	Accept the manual bank-verification workload	W3	Payout release in W7 needs verified bank accounts. Deferred, the first payout run either blocks or releases against unverified accounts, and only one of those is acceptable.	Roughly ten minutes of staff time per week at pilot scale: a ₹1 UPI transfer and a name match per account. Buying it out costs about ₹3 per verification.
K-4	Buy masked click-to-call, or ship the manifest reveal	W5	The manifest endpoint is built in W6 and its disclosure semantics are the reveal decision. Retrofitting a call bridge afterwards means changing the one endpoint in the system that releases an address, which is the endpoint you least want to change twice.	About ₹1 500 per month plus ₹0.60 per minute, and a written waiver against I-14 because it breaks the Rs 0 recurring rule. See R-04.
K-5	Set the COD shortfall write-off threshold and the signatory	W7	Settlement ships in W7. Without a threshold, the first SHORT remittance sits on the exception board with nobody empowered to clear it, and shortfalls accumulate as unreported losses.	Free to decide, and it costs whatever you set it to. The plan deliberately does not automate the write-off. See R-09.
K-6	Freeze the contract package	W8	Three portals compile against one published contract version. A breaking change in W9 means three repositories to re-release inside the hardening week, which is the week with the least room for it.	Free, and it costs flexibility: anything discovered in W9 is a forward-fix behind a feature flag rather than a contract change.
K-7	Buy automated verification or a BI layer, or neither	W8	Both are conveniences rather than blockers, which is exactly why the decision drifts. Deferred indefinitely is a legitimate answer here; what is not legitimate is planning as though someone will buy them and leaving the manual work unstaffed.	About ₹3 per bank verification. A hosted BI layer starts near \$85 per month, against eleven fixed dashboard tiles and masked CSV export that already exist.
K-8	Commit to the pilot cutover date	W9	The pilot has external participants with their own calendars. Deferred into W10, the pilot slips out of the programme entirely and the last week becomes an unbounded hardening week.	Free, and it is the decision that makes the cut list real: committing to the date means agreeing in advance that scope moves rather than the date.

None of these eight can be delegated to an engineer, and all eight have a week attached because a decision without a deadline is a decision that gets made by whatever happens first. Three of them — K-1, K-2 and K-8 — are the load-bearing ones, and they share a property worth noticing: each converts an open question into a constraint the team can plan against. A frozen supplier list is more useful to the team than a longer one, and a committed cutover date is more useful than a generous one, because both make the cut list in section 10.6 something the team can act on at a Friday gate without asking permission.

K-4 and K-7 are the two that spend money, and they are the only decisions in this programme that require a written waiver against the Rs 0 rule. Neither is a blocker. Both are worth revisiting after the pilot rather than before it, because the pilot is what turns "riders may need to call customers" into a countable number of failed deliveries, and a countable number is a much better basis for a ₹1 500 monthly commitment than an intuition.

10.8 MONEY

Rs 0 recurring, a shopping list with triggers, and what Rs 0 actually costs

RECURRING INFRASTRUCTURE

Rs 0

unchanged by three new front ends

AMORTISED MONTHLY TOTAL

₹67

the domain, already paid

NEW HOSTNAMES

6

admin, seller, fleet and the existing three

BUDGET LINES TRACKED

16

each with a free-tier ceiling and an alert

NAMED PAID GAPS

5

none of them a blocker

WAIVERS REQUIRED

0

unless K-4 or K-7 is taken

₹67 per month, unchanged from the pre-portal baseline. Three new front ends, one new package registry and three new hostnames add ₹0. Three new front ends, a package registry and three new hostnames add nothing, because Cloudflare Pages bills builds rather than sites and the portals are low-traffic by nature: tens of staff, tens of suppliers, hundreds of riders. AA-09 sections 9.1 and 9.2 hold the per-service arithmetic behind that claim across all 16 budget lines, including the polling calculation that puts the programme at roughly 1.53 million Cloud Run requests per month against a 2 million always-free ceiling. This section covers only the part a manager decides: what is not free, what it costs, and when to buy it.

THE GAP	WHAT SHIPS INSTEAD	REAL PRICE	OPERATIONAL TRIGGER TO BUY
Masked click-to-call between rider and customer	Time-boxed address reveal on the manifest, logged in pii_disclosures. Customer contact happens through a WhatsApp template send triggered by the rider, not a phone call.	Exotel or Kaleyra number masking, roughly ₹0.60/minute plus a ₹1 500/month platform fee.	Buy when more than five deliveries in a week fail because a rider could not reach a customer. Below that, the WhatsApp template send is doing its job. See R-04 and K-4.
Automated bank-account verification (penny drop)	A staff member transfers ₹1 by UPI and matches the holder name, then sets verified_at. Two suppliers and one fleet is roughly ten minutes of work per week.	Cashfree or Razorpay account validation, roughly ₹3 per verification.	Buy when more than about twenty accounts a week need verifying, at which point the staff time exceeds the fee. At two suppliers and one fleet this is nowhere near true. See K-3.
Real-time rider GPS tracking on a map for the customer	Status milestones only: HANDED_OVER, IN_TRANSIT, OUT_FOR_DELIVERY. The POD captures a single geo point, which is all the geofence needs.	Google Maps Platform beyond the \$200 monthly credit, plus a location ingest path.	Buy only if customers ask for it, and treat it as a Phase 2 programme with its own budget rather than a line item. Milestone statuses are what the mobile tracker already shows.
SMS delivery notifications at scale	WhatsApp template messages first, FCM push second, SMS only for OTP fallback.	2Factor.in or MSG91 at roughly ₹0.18 per transactional SMS plus DLT fees.	Buy when WhatsApp template delivery falls below about ninety per cent for order-critical notifications. Note DLT registration is a lead time, not just a price.
A hosted BI or reporting layer over the portal data	Eleven fixed dashboard tiles computed by one SQL query, cached 30 s at the edge, plus masked CSV export for anything deeper.	Metabase Cloud from roughly \$85/month, or a small always-on VM.	Buy when someone outside the team is asking the same data question weekly and the eleven fixed tiles plus masked CSV export cannot answer it. That is a Phase 2 signal, not a W10 one.

Every row above has a trigger rather than a date, and that is the intended discipline: these are purchases justified by an observed number, not by a roadmap slot. Two of them are already decisions in section 10.7 with a week attached, because they interact with code being written in that week. The other three are genuinely open and can stay open through the pilot.

What Rs 0 actually costs, stated honestly

The recurring bill is Rs 0 and the engineering cost of that is not zero. Three places in this programme exist only because a paid service was declined, and it is worth naming the price in days rather than letting Rs 0 read as free. The GitHub Actions cron driving a Postgres job_queue with FOR UPDATE SKIP LOCKED exists because Cloud Run scales to zero and a managed scheduler costs money; it is roughly three dev-days to build, it has a five-minute granularity floor, and it offers best-effort rather than exactly-once execution, which is why every deadline in the system is computed from a stored timestamp and evaluated lazily on read as well as by cron. The audit_log archival cron to R2 exists because Neon's free tier is 0.5 GB; that is another day and a half. The two caches that pull Neon compute-hours from 264 down to 154 exist because the free ceiling is 190; call that

two days. Six to seven dev-days of the 200 committed are spent buying Rs 0 rather than building features.

The second cost is manual work absorbed by staff, and it compounds with scale rather than staying flat. Bank verification is about ten minutes a week at pilot scale and becomes an hour a week at twenty suppliers. Commission invoices under cut C-3 are manual for two pilot suppliers and unmanageable at twenty. Service-area edits under cut C-8 are fine with three fleets and stated in the cut list itself as not scaling past ten. None of this is hidden, but it does mean the Rs 0 position is a pilot-scale position: it holds for the ten weeks and for the pilot, and the first real growth quarter will convert two or three of these gaps into purchases. Planning as though Rs 0 is permanent is the one way to be surprised by this.

The third cost is the one with no workaround, and it is not a cost of Rs 0 at all. Wallet cash-out is prohibited rather than deferred: paying money out of a stored balance makes the business a prepaid payment instrument issuer requiring RBI authorisation. No amount of budget changes that inside ten weeks. Wallet credit is spendable inside Atavishaala and refundable to the original payment instrument, and it never becomes cash.

10.9 HOW PROGRESS IS MEASURED

Four meetings, four numbers, and twenty things that have to be demonstrable

WHEN	LENGTH	MEETING	WHO	WHAT ACTUALLY HAPPENS IN IT
Monday 09:00	30 min	Week kick-off	All five	Read the week goal from AA-10 aloud. Each dev states their task IDs and the one thing that could block them. Cut-list status is stated even when nothing is cut.
Wednesday 18:00	20 min	Blocker sync	All five	Blockers only. No status. The kit release for the week ships at the end of this call, so FE-2 raises kit gaps here or waits a week.
Friday 16:00	45 min	Demo gate	All five + you	Working software against real data on the deployed hosts, not localhost. The week is either passed or explicitly failed; a failed gate moves work into the next week and triggers a cut-list decision.
Friday 17:00	15 min	Budget and threshold check	You + BE-1	Read the eight free-tier gauges. Any gauge above 70 % gets a named action before Monday.

Four meetings a week totalling under two hours, and only one of them involves you for its full length. The Friday demo gate is the one that matters, and it has a rule that is easy to state and easy to erode: the week is demonstrated against real data on the deployed hosts, not on a laptop, and the week is either passed or explicitly failed. There is no partially passed week. A merged pull request is not a passed week, because merged code that nobody has operated in a browser is code whose defects are still in front of you rather than behind you.

Four numbers get reported each Friday and they fit in a message. First, the gate: passed or failed, against the demonstrable end state for that week from the 10-row table in section 10.2. Second, the cut-list position: which C-numbers have been taken, stated even when the answer is none, because a silent cut list is how a programme discovers in week nine that four cuts were taken quietly. Third, the highest free-tier gauge as a percentage, from the Friday 17:00 check. Fourth, the count of risk indicators that fired this week, by ID, from the 12-row register in section 10.5.

ID	DEFINITION OF DONE — DEMONSTRABLE, NOT MERGED
D-01	A staff user can sign in with email, password and TOTP, and cannot sign in without TOTP.
D-02	Effective permissions resolve with DENY winning, and the resolution is unit-tested against a matrix.
D-03	Every mutating admin request appears in the audit log with a before/after diff.
D-04	No screen in any portal renders more than the last four characters of a masked field.
D-05	A product company can be invited, submit seven KYC documents, be verified and go ACTIVE.
D-06	A partner-drafted product passes moderation and appears in the mobile app catalog.
D-07	A multi-supplier order splits into shipments, and the app tracker shows a correct derived status.
D-08	A supplier can accept, pack, declare a shortfall, mark ready and hand over.
D-09	A delivery company can onboard, define service areas with a rate card, and add riders.
D-10	An offer broadcasts, expires, escalates through three waves and lands on the admin board.

ID	DEFINITION OF DONE — DEMONSTRABLE, NOT MERGED
D-11	A rider opens the manifest, the disclosure is registered, and delivery completes with OTP inside 200 m.
D-12	Three failed attempts force RTO automatically and the customer is refunded on RTO_DELIVERED.
D-13	A weekly payout run assembles, locks, releases with a UTR, and reconciles to the paisa.
D-14	A COD period is declared by a fleet and reconciled by finance, with variance computed.
D-15	A return with photo evidence is approved, picked, received by the supplier and refunded, with a compensating ledger entry.
D-16	An agronomist claims an advisory request from the admin portal and records a callback within the SLA.
D-17	All twelve CI gates pass on all five repositories.
D-18	M1..M8 apply cleanly to a restored production snapshot and the derived-status diff is zero rows.
D-19	The monthly bill is ₹67 and the eight free-tier alert thresholds are configured and firing on test.
D-20	Eleven runbooks are written and each has been rehearsed once by someone other than its author.

These 20 items are the programme-level gate, not the weekly one. Each is written as something a person can be shown in a browser, which is deliberate: "the permission matrix is implemented" is unfalsifiable, whereas D-02 — permissions resolve with DENY winning, unit-tested against a matrix — either happens in front of you or does not. Three of them are the ones to check personally rather than take on report, because they are the ones where a plausible-looking demonstration can hide a real gap: D-04, that no screen anywhere renders more than the last four characters of a masked field; D-07, that a multi-supplier order splits and the mobile tracker still shows a correct derived status; and D-18, that the migrations apply to a restored production snapshot with a zero-row status diff.

10.10 AFTER WEEK TEN

Seven phase-2 candidates, and which ones the pilot will make urgent

#	CANDIDATE	HOW URGENT	SIZE	THE REASONING
1	Saved views and CSV export (cut C-2)	Urgent within a month of go-live	3 dev-days	Ops will hit this first because it is the thing they do twenty times a day. It is a cut, not a new feature, so it is cheap and pre-scoped.
2	Self-billed commission invoice PDF (cut C-3)	Urgent at roughly five suppliers	2 dev-days	Manual invoicing works for two pilot suppliers and stops working quickly. Note AA-06 section 6.9 shows screen SP-22 promising this PDF, which C-3 removes; that inconsistency is flagged and the screen ships without the download until this lands.
3	Vehicle registry (cut C-7)	Urgent only if heavy equipment moves to third-party fleets	3 dev-days	Blocks nothing while heavy equipment routes to ATV-SELF. Becomes a blocker the week a fleet asks to carry it.
4	Service-area approval workflow (cut C-8)	Urgent past about ten fleets	3 dev-days	Admin editing fleet areas directly is fine at three fleets and is stated in the cut list itself as not scaling past ten.
5	Custom roles (cut C-4)	Can wait a quarter	4 dev-days	The nine system roles cover every seat that can be named today. This becomes real when a customer wants a seat that sees payouts but not PII, which the matrix already supports but the editor does not expose.
6	A real scheduler to replace the cron queue	Can wait a quarter, unless a deadline is missed materially	4 dev-days plus a bill	The five-minute granularity floor and best-effort execution are already worked around by lazy evaluation on read. Replace it when a business deadline actually depends on sub-five-minute timing, not before.
7	A public partner or fleet API	Not urgent; wait for a named integrator	unscoped	The webhook_deliveries table exists from M8 and stays empty. Building a public API before a partner with a system asks for it is the most expensive way to guess.

The ordering is by what the pilot will make urgent, not by what is interesting to build, and the first four are all cuts rather than new work. That is the payoff of writing the cut list before the pressure arrived: the Phase 2 backlog is mostly pre-scoped, pre-estimated items whose consequences were already written down and

accepted, so the first post-pilot planning session starts from a list rather than from a discussion. Items one through four total eleven dev-days, which is under a fortnight of one engineer.

Two things on this list should be resisted rather than scheduled. The scheduler replacement is the kind of work engineers want to do and that a manager should require a missed business deadline to justify, because the cron queue is ugly and adequate. The public API should be resisted until a specific named partner with a specific system asks, because every guess about an integration contract becomes a compatibility obligation. The one item genuinely absent from this list is a native rider application, and its absence is a judgement: if the pilot shows riders cannot work in a mobile browser at the doorstep, that is a programme of its own with its own charter, not a Phase 2 line item.

10.11 HOW TO KILL THIS PROGRAMME

Three conditions under which stopping is the right call, and what survives if you do

Every plan of this size should say how it ends badly, because the alternative is a programme that is extended three times by people who each individually could not see a reason to stop. The three conditions below are not failure modes of the engineering; the code will very likely work. They are failure modes of the thesis — that suppliers want self-service, that third-party fleets want this commission, and that automation removes more staff work than the Rs 0 workarounds create. Each has a measurement, a week by which it is visible, and a salvage position, because the salvage is what makes stopping a decision rather than a defeat.

#	KILL CONDITION	HOW IT IS MEASURED	VISIBLE BY
1	The pilot suppliers do not use the portal	Fewer than half of the shipments in a week are accepted, packed and handed over by a supplier user rather than by staff acting on their behalf; and fewer than five products are drafted by a supplier in the first fortnight after go-live.	W10 plus two weeks
2	Third-party fleets will not accept the commission economics	Two or more fleets reach a rate-card conversation and decline, or fleet acceptance rate on broadcast offers stays below 40 % with median accept latency above eight minutes once real volume exists.	W8 for the conversations, W10 plus two weeks for the acceptance rate
3	The manual workaround load exceeds the staff time saved	Staff hours per week spent on bank verification, manual invoicing, service-area edits, COD chasing and exception-board clearing, compared against the hours those same people spent on manual onboarding and shipment coordination before W1.	W10 plus four weeks, and it needs the before-number recorded in W1

Condition three has a trap in it worth naming now: the comparison is impossible unless somebody writes down, in week one, how many hours a week staff currently spend onboarding suppliers, cataloguing products and coordinating multi-supplier orders by hand. That number is the entire business case from section 10.1, and it is nobody's task in AA-11. Record it in week one or accept that the programme cannot be evaluated on the thing it was justified by.

Condition two is the one most likely to fire, and it is also the one where stopping is least likely to be correct — which is exactly why it needs a written measurement rather than a judgement in the moment. If fleets decline the commission, the fallback in R-02 is already built: ATV-SELF operates the fleet portal, staff deliver, and every code path is exercised. That is not the marketplace, but it is a working delivery operation with offers, POD, geofencing and a fleet ledger, and it is a better base to negotiate a rate card from than a slide deck. Stopping is correct only if the economics are wrong at any commission a fleet would accept, which is a pricing conclusion and not a delivery one.

WHAT SURVIVES A STOP, WHATEVER THE REASON	WHY IT KEEPS ITS VALUE
The staff identity system, the 88-permission matrix, TOTP and the audit log	The founder console is replaced regardless. Nine roles, DENY-wins resolution and a before/after diff on every mutation are worth having even if no external tenant ever signs in.
The last-four-only masking regime and the pii_disclosures register	This is a compliance posture, not a feature. It applies to the existing mobile app and to every future surface, and retrofitting it later costs an audit of every query rather than a Prisma extension.
The shipment split and the derived order status	Multi-supplier orders are a data-model correction that the platform needed anyway. It is the change that makes staff coordination tractable even with no supplier portal at all.
The append-only ledgers in integer paise	Commission and delivery-fee accrual is answerable from data. This survives any decision about who does the delivering.
The agronomist console and the two-hour advisory SLA	Crop Advisor is independent of the marketplace thesis entirely. It ships in week eight and is unaffected by all three kill conditions.

WHAT SURVIVES A STOP, WHATEVER THE REASON	WHY IT KEEPS ITS VALUE
The portal kit and the five-repository split	Any future surface — a rider app, a phase-2 portal, an internal tool — starts from a published component kit, a contract package and a CI pipeline that already work.

The list above is why the ordering in section 10.2 matters beyond engineering hygiene. Read-only before write, tenancy before features, money last means the work that survives a stop is front-loaded, and the work that depends on an unproven thesis is at the back. A stop decision in week eight leaves you with identity, permissions, audit, masking, the shipment split and the moderation pipeline — six things worth keeping. The same decision under a plan that had built the fleet marketplace first would leave you with an offer allocator and no audit log.

One closing note on how to use this section. Do not consult it for the first time in the week you are worried. Read the three measurements at the week-one kick-off, get the week-one baseline number recorded, and state at each Friday gate whether any of the three is trending. A kill condition that is checked weekly is a management instrument. One that is remembered in week twelve is a post-mortem.

THE ONE CONDITION THAT IS NOT ON THIS LIST

Ten weeks with four engineers being optimistic is not a kill condition. It is R-12, rated high probability, and its response is the ordered cut list in section 10.6 followed by moving the cutover date. A programme that runs to thirteen weeks and delivers three working portals has succeeded at the thing it was funded for. Confusing a schedule overrun with a failed thesis is the most likely way for this plan to be stopped for the wrong reason.